

X Data Analysis using API

Nguyen Duc Thanh Vinh - 23229582
Nguyen The Hung - 23229586
Nguyen Phuong Huyen - 23229585

Group 1

Table of Contents

1. Data collection
2. Data preprocessing
3. Hypothesis

Data Collection

```
#Set up your Bearer Token
bearer_token = "AAAAAAAAAAAAAAAMgV4wEAAAAA09spRs1Kr%2BZUrWP0JzmFicMNxYo%3DZDzkxtTIzNn5A7RhwdfBPq1XESws44pYLqDwGZwZ7GeVKVDL6z"
headers = {"Authorization": f"Bearer {bearer_token}"}


#Get user ID from username (no @ symbol)
username = "Google"
user_url = f"https://api.twitter.com/2/users/by/username/{username}"
user_resp = requests.get(user_url, headers=headers).json()
#The .json() part at the end of this line converts the response from the API into a Python dictionary
```

```
# Safety check before accessing 'data'
if "data" in user_resp:
    user_id = user_resp["data"]["id"]
# This line accesses the value associated with the key "id", which is inside the key "data" in t

    #Fetch recent tweets
    tweet_url = f"https://api.twitter.com/2/users/{user_id}/tweets"
    params = {
        "max_results": 100,
        "tweet.fields": "created_at,public_metrics,text"
    }
    #Catch the value from the key "tweet.fields" , asking for tweet's creation time (created_at
}
    tweet_resp = requests.get(tweet_url, headers=headers, params=params).json()

    if "data" in tweet_resp:
        #Load tweets into DataFrame
        df = pd.json_normalize(tweet_resp["data"])
        if "public_metrics.like_count" in df.columns:
            df["created_at"] = pd.to_datetime(df["created_at"])
            df["likes"] = df["public_metrics.like_count"]
            df["retweets"] = df["public_metrics.retweet_count"]
            df["hour"] = df["created_at"].dt.hour
                # Plot average likes by hour
                avg_likes_by_hour = df.groupby("hour")["likes"].mean()
                avg_likes_by_hour.plot(kind="line", marker="o", color="darkblue")
                plt.title(f"Average Likes by Hour of Day ({username})")
                plt.xlabel("Time of Day")
                plt.ylabel("Average Likes")
                plt.grid(True)
                plt.tight_layout()
                plt.show()
            else:
                print("Expected metrics not found in tweet data. Columns available:")
                print(df.columns)
        else:
            print("No tweet data returned. Full response:")
            print(tweet_resp)
        else:
            print("User not found or token issue. Full response:")
            print(user_resp)
            # Save to CSV in your Google Drive
            df.to_csv('/content/drive/MyDrive/google_tweets.csv', index=False)
```

Data checking, save to csv ***limit 100 tweets per time**



Data Processing

- rename ‘hour’ → ‘posting_time’
- drop two columns: ‘likes’ and ‘retweets’

#	Column
---	---
0	text
1	edit_history_tweet_ids
2	created_at
3	id
4	public_metrics.retweet_count
5	public_metrics.reply_count
6	public_metrics.like_count
7	public_metrics.quote_count
8	public_metrics.bookmark_count
9	public_metrics.impression_count
10	likes
11	retweets
12	hour
13	article.title

before

#	Column
---	---
0	text
1	edit_history_tweet_ids
2	created_at
3	id
4	public_metrics.retweet_count
5	public_metrics.reply_count
6	public_metrics.like_count
7	public_metrics.quote_count
8	public_metrics.bookmark_count
9	public_metrics.impression_count
10	posting_time

after

Data Processing

removed tweets with 'public_metrics.impression_count' == 0

```
# remove impression_count = 0
tweet_data = tweet_data[tweet_data['public_metrics.impression_count'] > 0]
```

**calculate the engagement rate
with the following formula:**

$$\text{engagement rate} = \frac{\text{total interaction}}{\text{impression count}}$$

**with total interaction = sum of reply,
like, retweet, quote, bookmark**

```
# Calculate engagement_rate
tweet_data['engagement_rate'] = (
    (tweet_data['public_metrics.reply_count'] +
     tweet_data['public_metrics.retweet_count'] +
     tweet_data['public_metrics.like_count'] +
     tweet_data['public_metrics.quote_count'] +
     tweet_data['public_metrics.bookmark_count'])
    / tweet_data['public_metrics.impression_count'])
```

Hypothesis 1

Research question: How does the posting time (hour of the day) affect the engagement of tweets?

Null Hypothesis (H_0): The posting time (hour of the day) has no significant effect on tweet engagement.

Alternative Hypothesis (H_1): The posting time (hour of the day) significantly improves tweet engagement.

1. group posting_time into four groups: Morning, Afternoon, Evening, Night

```
# Group posting_time
def categorize_hour(hour):
    if 5 <= hour <= 11:
        return 'Morning'
    elif 12 <= hour <= 17:
        return 'Afternoon'
    elif 18 <= hour <= 22:
        return 'Evening'
    else:
        return 'Night'

tweet_data['hour_group'] = tweet_data['posting_time'].apply(categorize_hour)
```

3. remove outliers using IQR method (1.5 x IQR rule)

```
def remove_outliers_by_group(df, group_col, value_col):
    cleaned_df = pd.DataFrame()

    for group, data in df.groupby(group_col):
        Q1 = data[value_col].quantile(0.25)
        Q3 = data[value_col].quantile(0.75)
        IQR = Q3 - Q1

        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR


        filtered = data[(data[value_col] >= lower_bound) & (data[value_col] <= upper_bound)]
        cleaned_df = pd.concat([cleaned_df, filtered])

    return cleaned_df
tweet_clean = remove_outliers_by_group(tweet_data, 'hour_group', 'engagement_rate')

print("Before:", len(tweet_data))
print("After:", len(tweet_clean))
```

Hypothesis 1


2. create box plot to check if there are outliers



Hypothesis 1

4. check skewness

	hour_group	original_skew	log_skew
0	Afternoon	1.292876	1.270862
1	Evening	1.283973	1.259993
2	Morning	1.358648	1.320552
3	Night	0.967488	0.949055



after remove outliers and check skewness again --> distribution each groups is not normal --> we decide to use Kruskal-Wallis test (non-parametric tests)

```
stat, p = kruskal(*[group['engagement_rate'].values for name, group in tweet_data.groupby('hour_group')])  
print(f"Kruskal-Wallis test: H={stat:.3f}, p={p:.4f}")  
  
alpha = 0.05  
if p < alpha:  
    print("Since p-value < 0.05, we reject the null hypothesis.")  
    print("There is a statistically significant difference in tweet engagement across posting time groups.")  
else:  
    print("Since p-value ≥ 0.05, we fail to reject the null hypothesis.")  
    print("There is no statistically significant difference in tweet engagement across posting time groups.")
```

Conclusion for hypothesis 1

after running code of Kruskal-Wallis test, this is the result:

Kruskal-Wallis test: $H=5.631$, $p=0.1310$

Conclusion for Hypothesis 1: How does the posting time (hour of the day) affect the engagement of tweets?

Since the p_value > 0.05, we fail to reject the null hypothesis

→ This indicates that there is no statistically significant difference in tweet engagement across different posting time groups

Hypothesis 2

Research question: How does the number of bookmarks affect the number of likes on tweets?

Null Hypothesis (H_0): There is no significant difference in the number of likes between tweets with higher and lower bookmark counts.

Alternative Hypothesis (H_1): Tweets with higher bookmark counts have significantly more likes than those with lower bookmark counts.


Hypothesis 2

1. divide high and low bookmarks according to median:

```
median_bookmarks = tweet_data['public_metrics.bookmark_count'].median()  
tweet_clean['bookmark_group'] = np.where(  
    tweet_clean['public_metrics.bookmark_count'] >= median_bookmarks, 'High', 'Low'  
)
```

2. create a box plot

tweets with high bookmark counts tend to receive significantly more likes, suggesting a positive relationship between bookmarking and user appreciation.



Hypothesis 2

3. The Mann-Whitney U test was used to compare the distributions of engagement levels between the two groups.

```
# Mann-Whitney U test
high = tweet_clean[tweet_clean['bookmark_group'] == 'High']['public_metrics.like_count']
low = tweet_clean[tweet_clean['bookmark_group'] == 'Low']['public_metrics.like_count']

stat, p = mannwhitneyu(high, low, alternative='greater')

print(f"Mann-Whitney U statistic: {stat:.4f}, p-value: {p:.4e}")

if p < 0.05:
    print("Tweets with higher bookmarks tend to have more likes (significant difference).")
else:
    print("No significant difference in likes between high and low bookmark groups.")
```

Conclusion for hypothesis 2

After running code of Mann - Whitney U test, this is the result:

Mann-Whitney U statistic: 93334.000, p-value: 4.2782e-97

Conclusion for Hypothesis 2: How does the number of bookmarks affect the number of likes on tweets?

Since the p_value < 0.05, we reject the null hypothesis

→ Tweets with higher bookmarks tend to have more likes (significant difference)

Hypothesis 3

Research question: How does tweet length affect user engagement on Twitter?

- Null Hypothesis (H_0): Longer tweets do not receive more likes than shorter tweets.
- Alternative Hypothesis (H_1): Longer tweets tend to receive more likes than shorter tweets.


Hypothesis 3

1. add a new column representing of each tweet (in characters)

```
df['tweet_length'] = df['text'].astype(str).apply(len)
```

2. create a scatter plot

**tweets with moderate length
tend to get higher likes
→ use one-tailed test**



Hypothesis 3

3. check the correlation between tweet length and likes/retweets

```
likes_corr = df['tweet_length'].corr(df['likes'])  
print(f"Correlation with Likes: {likes_corr:.3f}")
```

Correlation with likes = 0.237

→ Longer tweets slightly tend to receive more likes, but the relationship is weak.

4. Split tweets into two groups

- Short tweets: shorter than or equal to mean
- Long tweets: longer than mean

```
df['tweet_length_mean']=df['tweet_length'].mean()  
short_tweets = df[df['tweet_length'] <= df['tweet_length_mean']]  
long_tweets = df[df['tweet_length'] > df['tweet_length_mean']]
```

Hypothesis 3

5. Test whether tweet length affects the number of likes by using a one-tailed test

```
from scipy.stats import ttest_ind

t_likes = ttest_ind(long_tweets['likes'], short_tweets['likes'], equal_var=False, alternative ='greater')

print(f"T-test for Likes → t-statistic: {t_likes.statistic:.5f}, p-value: {t_likes.pvalue:.5f}")
```

Conclusion for hypothesis 3

After running the code of t-test, this is the result:

T-test for Likes → t-statistic: 3.62649, p-value: 0.00016

Since the p-value is less than 0.05, we reject the null hypothesis.

→ **Longer tweets tend to receive more likes and on average than shorter ones, suggesting that tweet length positively influences user engagement.**

References

- Kruskal, W. H., & Wallis, W. A. (1952). Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260), 583–621.
- Mann, H. B., & Whitney, D. R. (1947). On a test of whether one of two random variables is stochastically larger than the other. *The Annals of Mathematical Statistics*, 18, 50–60.
- Yuen, K. K. (1974). The two-sample trimmed t for unequal population variances. *Biometrika*, 61(1), 165–170.

Thank You